

4.0 Ensemble Learning Overview

Jonathan De Los Santos

2/13/2021

Contents

1	Ensemble Forms	1
2	Bagging	1
2.1	Bagging in R	2
3	Boosting	5
3.1	Bagging vs Boosting	6
3.2	Boosting in R Overview	6
3.3	Evaluating Boosting Model	8
4	Extensions	9
4.1	Random Forest	10

1 Ensemble Forms

- Bagging and random forest
 - Aggregate pre-determined weak learners
- Boosting
 - Gradually connect weak learners

2 Bagging

- Effectively reduces the variance of an individual base learner
 - Improves prediction accuracy for high variance
 - The trees are not independent, called “tree correlation”
- Doesn’t always improve upon an individual base learner
- Works for unstable and high variance learners
 - Decision trees and kNN
- Categorical Data

- Most predicted values selected as final forecasts
- Continuous Data
 - Average of predicted values determines final forecasts

2.1 Bagging in R

```
library(adabag)
library(rpart)
library(rpart.plot)
library(caret)

set.seed(123)
```

```
idx<-sample(1:150,100)
train<- iris[idx,]
test<- iris[-idx,]

# Single Tree
iris.single <- rpart(Species ~ ., data = train, method= 'class')

# Bagging with 5 trees
iris.bagging <- bagging(Species~ ., data = train, mfinal =5,
                       control = rpart.control(maxdepth=5, minsplit=5))


iris.bagging$importance
```

```
## Petal.Length  Petal.Width Sepal.Length  Sepal.Width
##      69.14617      30.85383      0.00000      0.00000
```

2.1.1 Bagging Plot

```
rpart.plot(iris.bagging$trees[[1]])
```


```
## Warning: Cannot retrieve the data used to build the model (so cannot determine roundint and is.binary)
## To silence this warning:
##   Call rpart.plot with roundint=FALSE,
##   or rebuild the rpart model with model=TRUE.
```



2.1.2 Bagging Barplot

```
barplot(iris.bagging$importance[order(iris.bagging$importance, decreasing = TRUE)], ylim = c(0, 100), m
```

Variable Relative Importance



Evaluating Bagging Model

```
library(e1071)
pred <- predict (iris.single, test, type = "class")
confusionMatrix(pred, test$Species)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  setosa versicolor virginica
##   setosa      16          0          0
##   versicolor   0          20         1
##   virginica    0           1        12
##
## Overall Statistics
##
##              Accuracy : 0.96
##              95% CI : (0.8629, 0.9951)
##   No Information Rate : 0.42
##   P-Value [Acc > NIR] : 3.498e-16
##
##              Kappa : 0.9388
##
##   McNemar's Test P-Value : NA
##
## Statistics by Class:
```

```
##
##               Class: setosa Class: versicolor Class: virginica
## Sensitivity           1.00           0.9524           0.9231
## Specificity           1.00           0.9655           0.9730
## Pos Pred Value        1.00           0.9524           0.9231
## Neg Pred Value        1.00           0.9655           0.9730
## Prevalence            0.32           0.4200           0.2600
## Detection Rate        0.32           0.4000           0.2400
## Detection Prevalence  0.32           0.4200           0.2600
## Balanced Accuracy      1.00           0.9589           0.9480
```

```
pred2 <- predict(iris.bagging, test, type = "class")

# Convert pred2 class to factor before building confusion matrix
confusionMatrix(factor(pred2$class), test$Species)
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction  setosa versicolor virginica
##   setosa      16          0          0
##   versicolor   0         19          0
##   virginica    0          2         13
##
## Overall Statistics
##
##               Accuracy : 0.96
##               95% CI : (0.8629, 0.9951)
##   No Information Rate : 0.42
##   P-Value [Acc > NIR] : 3.498e-16
##
##               Kappa : 0.9394
##
##   McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: setosa Class: versicolor Class: virginica
## Sensitivity           1.00           0.9048           1.0000
## Specificity           1.00           1.0000           0.9459
## Pos Pred Value        1.00           1.0000           0.8667
## Neg Pred Value        1.00           0.9355           1.0000
## Prevalence            0.32           0.4200           0.2600
## Detection Rate        0.32           0.3800           0.2600
## Detection Prevalence  0.32           0.3800           0.3000
## Balanced Accuracy      1.00           0.9524           0.9730
```

3 Boosting

Boosting

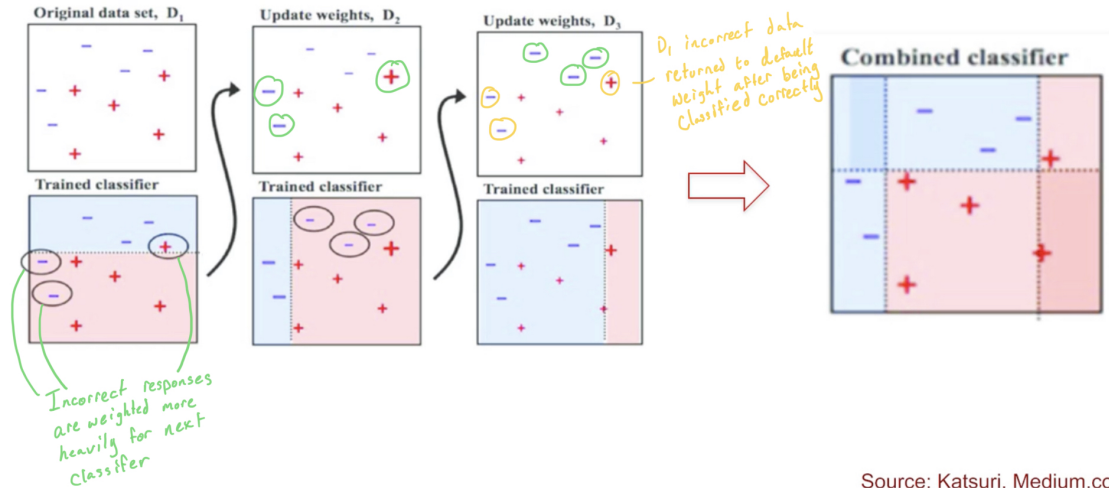


Figure 1: Boosting

3.1 Bagging vs Boosting

Bagging: - Taught in parallel - Decreases variance - More errors - Appropriate when overfit is a problem
 Boosting: - Taught sequentially - Decreases bias - Slow and may overfit - Appropriate when low performance is a problem

3.2 Boosting in R Overview

Uses similar setup as bagging with a few differences

```
# Single tree setup
iris.single <- rpart(Species ~ ., data = train, method= 'class')

# Boosting with 5 trees
# Note change from "bagging" to "boosting" model, otherwise same as bagging
iris.boosting<- boosting(Species~ ., data = train, mfinal =5,
                        control = rpart.control(maxdepth=5, minsplit=5))

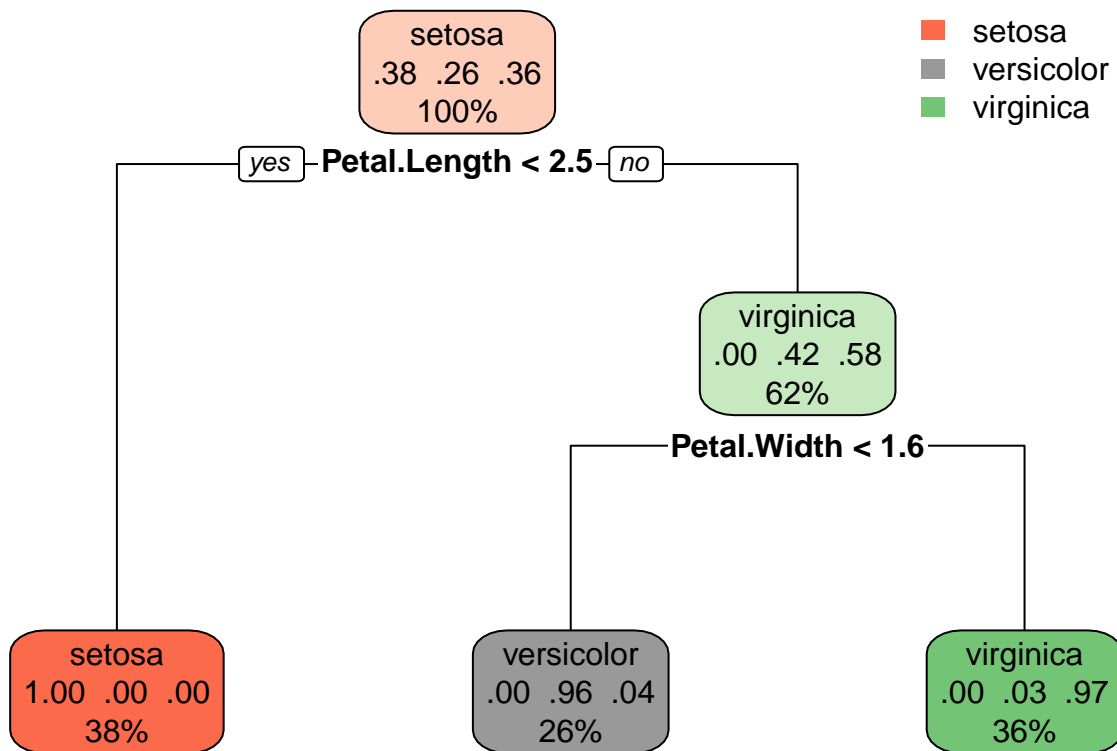
iris.boosting$importance
```

```
## Petal.Length Petal.Width Sepal.Length Sepal.Width
##      58.238326      36.321655       5.024207       0.415812
```

3.2.1 Boosting Plot

```
rpart.plot(iris.boosting$trees[[1]])
```


```
## Warning: Cannot retrieve the data used to build the model (so cannot determine roundint and is.binary)
## To silence this warning:
##   Call rpart.plot with roundint=FALSE,
##   or rebuild the rpart model with model=TRUE.
```



3.2.2 Boosting Barplot

Note the extremely important difference from bagging: we made it green

```
barplot(iris.boosting$importance[order(iris.boosting$importance, decreasing = TRUE)], ylim = c(0, 100),
```



3.3 Evaluating Boosting Model

```
library(e1071)
pred3 <- predict(iris.single, test, type = "class")
confusionMatrix(pred3, test$Species)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  setosa versicolor virginica
## setosa      16          0          0
## versicolor   0         20          1
## virginica    0          1         12
##
## Overall Statistics
##
##              Accuracy : 0.96
##              95% CI : (0.8629, 0.9951)
## No Information Rate : 0.42
## P-Value [Acc > NIR] : 3.498e-16
##
```



```
##                      Kappa : 0.9388
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##                      Class: setosa Class: versicolor Class: virginica
## Sensitivity           1.00          0.9524          0.9231
## Specificity           1.00          0.9655          0.9730
## Pos Pred Value        1.00          0.9524          0.9231
## Neg Pred Value        1.00          0.9655          0.9730
## Prevalence            0.32          0.4200          0.2600
## Detection Rate        0.32          0.4000          0.2400
## Detection Prevalence  0.32          0.4200          0.2600
## Balanced Accuracy      1.00          0.9589          0.9480
```

```
pred4 <- predict(iris.boosting, test, type = "class")
confusionMatrix(factor(pred4$class), test$Species)
```

```
## Confusion Matrix and Statistics
##
##                      Reference
## Prediction  setosa versicolor virginica
## setosa      16          0          0
## versicolor  0          19          0
## virginica   0           2         13
##
## Overall Statistics
##
##                      Accuracy : 0.96
##                      95% CI : (0.8629, 0.9951)
## No Information Rate : 0.42
## P-Value [Acc > NIR] : 3.498e-16
##
##                      Kappa : 0.9394
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##                      Class: setosa Class: versicolor Class: virginica
## Sensitivity           1.00          0.9048          1.0000
## Specificity           1.00          1.0000          0.9459
## Pos Pred Value        1.00          1.0000          0.8667
## Neg Pred Value        1.00          0.9355          1.0000
## Prevalence            0.32          0.4200          0.2600
## Detection Rate        0.32          0.3800          0.2600
## Detection Prevalence  0.32          0.3800          0.3000
## Balanced Accuracy      1.00          0.9524          0.9730
```

4 Extensions

These ensemble forms can be extended with further models:

Bagging \Rightarrow Random Forests Boosting \Rightarrow Gradient Boosting

4.1 Random Forest

4.1.1 Algorithm Structure

```
Given a training data set
Select number tree to build (n_trees)
for i=1 to n_tree do
  Generate a bootstrap sample
  Grow tree
    for
      | Select m_try variables at random from all p variables
      | Pick the best variable
      | Split the node into two child nodes end
    end
  Use tree model stopping criteria to determine when a tree is complete
end
```

4.1.2 Hyperparameters

- The number of trees in the forest
 - Start with $p \times 10$ trees and adjust as necessary
- For each split: `m_try`
 - The number of features to consider at any given split
 - Rule of thumb
 - * Classification: $m = \sqrt{p}$
 - * Regression: $m = p/3$
- Majority voting

4.1.3 Random Forests in R

- `nodesize` = 3 Number of nodes at the end of the tree
- `mtry` = 2 Number of variables

The confusion table shows us the error rate of each category

```
#install.packages("randomForest")
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 4.0.2
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
##
##     margin

# Create training data
train <- sample(1:150, 100)

# Build model
myForest <- randomForest (Species ~., nodesize = 3, mtry = 2, ntree = 15, data = iris)

# Observe confusion table
myForest$confusion
```

```
##           setosa versicolor virginica class.error
## setosa         50           0           0         0.00
## versicolor      0          47           3         0.06
## virginica       0           6          44         0.12
```

4.1.4 Evaluating Random Forest

```
# Load with test data
predict(myForest, newdata = iris[-train, ])
```

```
##           1           2           5           6           9          11          13
## setosa setosa setosa setosa setosa setosa setosa
## 14      25      26      28      29      33      35
## setosa setosa setosa setosa setosa setosa setosa
## 41      42      43      46      50      55      57
## setosa setosa setosa setosa setosa versicolor versicolor
## 60      61      64      65      68      71      75
## versicolor versicolor versicolor versicolor versicolor versicolor versicolor
## 76      80      82      83      84      88      89
## versicolor versicolor versicolor versicolor versicolor versicolor versicolor
## 94      95      98     101     102     115     117
## versicolor versicolor versicolor virginica virginica virginica virginica
## 124     126     134     136     137     141     148
## virginica virginica virginica virginica virginica virginica virginica
## 150
## virginica
## Levels: setosa versicolor virginica
```

```
# Create confusion matrix
tt <- table(iris$Species[-train], predict(myForest, iris[-train,]))
```

Calculate success and error rate:

```
# Success Rate
sum(tt[row(tt) == col(tt)]) / sum(tt)
```

```
## [1] 1
```

```
# Error Rate (1 - success rate)
1 - sum(tt[row(tt) == col(tt)]) / sum(tt)
```

```
## [1] 0
```

4.1.5 Evaluation Plotting

```
library(ggplot2)
test <- iris[-train,]
test$pred <- predict(myForest, iris [-train,])

ggplot(test, aes(Species, pred, color = Species)) +
  geom_jitter(width = 0.2, height = 0.1, size = 2) +
  labs(title = "Confusion Matrix",
       subtitle = "Predicted vs. Observed from Iris Dataset",
       y = "Predicted", x = "Truth")
```

